

**Asynchronous post commit triggers –
Replication to SPL routine –
Streaming analytics**

Agenda

- ▶ Use cases
- ▶ How it works
- ▶ Prerequisites
- ▶ Demo

Use cases

- ▶ Realtime Streaming analytics on OLTP data
 - MIN,MAX,AVG,SUM for a group of records.
 - Example: Per store sales reports
 - Realtime leaderboard calculation for an online game
 - Build materialized views
- ▶ Data transformations
 - Add additional fields ---like store id-- while replicating data to central server.
- ▶ Update external systems like graph database, Hadoop, Spark, Queuing services ...
- ▶ “No Key” data replication support
 - Replicate data for tables that do not have primary key, unique index or ER key.

How it works? (1)

- ▶ Solution is based on Enterprise Replication
- ▶ As part of transaction replay, instead of applying data to target table, user defined stored procedure gets fired for insert, update and delete operations.
- ▶ Works with loopback replication
 - Source and target participant can be defined on the same table-- acts likes a post commit asynchronous trigger
 - Source and target tables can be on the same server instance either in the same database or in different database
- ▶ Target table can be on a different Enterprise Replication server instance.
- ▶ Target table is used for parsing replicated row and extracting column values– data will not be applied to target table
- ▶ User can specify where clause filter to fire SPL on specific dataset!

How it works?(2)

- ▶ Data is staged in ER queues for asynchronous processing
- ▶ Source server id, transaction id, transaction commit time and operation type(I/U/D) are passed in as argument values to the SPL routine along with user data.
- ▶ For update operation, both before and after image of the column values are passed in as arguments to SPL routine.
- ▶ SPL routine execution can be configured to be invoked as user informix or table owner.

New options to 'cdr define replicate' command

- ▶ `--splname=<spl routine name>`
 - Stored procedure routine name to apply data to. SPL routine must exist at all participants
 - Input arguments: operation type, source id, txnid, before image of the row column list, after image of the row column list
- ▶ `--jsonsplname=<spl routine name>`
 - Stored procedure routine name to apply data to. SPL routine must exist at all participants
 - Input arguments: json document
 - `--jsonsplname` option expects input arguments to `splname` routine to be a json datatype. With json document as input to SPL routine, same SPL routine can be used for registering 'replication to SPL' replicate definition on multiple tables. For certain use cases --like queueing data to message queues -- this makes developer job a lot easier.
 - `--jsonsplname` option is mutually exclusive to `--splname` option.
- ▶ `--cascaderepl=y|n` enable cascade replication
 - Required if replication to SPL needs to be executed for the data applied through Enterprise Replication

--splname option stored procedure argument list

- ▶ Optype char(1) – operation type. Values include
 - I – Insert
 - U – Update
 - D – Delete
- ▶ Soucre_id integer – Source server id. Same as group id.
- ▶ Committime integer – Transaction commit time.
- ▶ Txnid bigint – Transaction id.
- ▶ Before value column list.
- ▶ After value column list.
 - Note: Column list for SPL routine extracted from select statement projection list

--jsonsplname option SPL routine json argument

Attribute name	Description
operation	Operation type: Insert/Delete/Update
table	Table name
owner	Table owner
database	Database name
txnid	8 byte unique id. Higher order 4 bytes: commit work log id, lower order 4 bytes: commit work log position.
commit_time	Transaction commit time for the event data.
rowdata	Row data in JSON document format. Data is returned in column name as key and column data as value.
before_rowdata	Before row data for “update” operation.

Example document format :

```
{“operation”:“insert”,“table”:“creditcardtxns”,“owner”:“informix”,“database”:“creditdb”,“txnid”:2250573177224,“commit_time”:1488243530,“rowdata”:{“uid”:22,“cardid”:“6666-6666-6666-6666”,“carddata”:{“Merchant”:“Sams Club”,“Amount”:200,“Date”:2017-05-01T10:35:10.000Z } }}
```

```
{“operation”:“update”,“table”:“creditcardtxns”,“owner”:“informix”,“database”:“creditdb”,“txnid”:2250573308360,“commit_time”:1488243832,“rowdata”:{uid:21,cardid:“7777-7777-7777-7777”,“carddata”:{“Merchant”:“Sams Club”,“Amount”:200,“Date”:“25-Jan-2017 16:15” } },“before_rowdata”:{“uid”:21,“cardid”:“6666-6666-6666-6666”,“carddata”:{“Merchant”:“Sams Club”,“Amount”:200,“Date”:2017-05-01T10:35:10.000Z } }}
```

```
{“operation”:“delete”,“table”:“creditcardtxns”,“owner”:“informix”,“database”:“creditdb”,“txnid”:2250573287760,“commit_time”:1488243797,“rowdata”:{“uid”:22,“cardid”:“6666-6666-6666-6666”,“carddata”:{“Merchant”:“Sams Club”,“Amount”:200,“Date”:2017-05-01T13:35:06.000Z } }}
```


Asynchronous post commit trigger support

- ▶ Define loopback replication server
- ▶ Create 'replication to SPL' type replicate with same "database and table" information for both source and target participants. Loopback server group name shall be specified with target participant definition:
- ▶ Example:
 - `cdr define repl rep1 -C always -S row -M g_cdr_utm_nag_1 -A -R --splname=logger4repl2spl "test@g_mygroup:informix.t1" "select * from t1" "test@g_loopback:informix.t1" "select * from t1"`
 - Note: g_mygroup is the local server ER group, and g_loopback is the pseudo ER server group.

Prerequisites

- ▶ Logging must be enabled for the tables
 - Does not work for RAW tables, and non-logged databases
- ▶ Cannot mix participant definitions that include table as the target, and SPL routine as the target.
 - Replicate definition marked with "repl2spl" attribute only support SPL routine as the target to apply data.
- ▶ Even though data is applied to stored procedure routine, target table definition must exist.
- ▶ Out-of-row data datatypes like text,byte,blob,clob aren't supported.

Use case 1 – Build staging table for data changes (--splname example)

```
create database test with log;
```

```
create table t1 (c1 int , c2 int);
```

```
create table t1log (optype int, srcdir int, committime int, txnid bigint, c1bef int, c2bef int, c1 int, c2 int);
```

```
create procedure logger4repl2spl (opType char(1), srcid integer, committime integer, txnid bigint, c1bef integer, c2bef int, c1 integer, c2 int)
```

```
    insert into t1log values (opType, srcid, committime, txnid, c1bef, c2bef, c1, c2);
```

```
end procedure;
```

```
$ cdr define repl rep1 -C always -S row -M g_cdr_utm_nag_1 -A -R --splname=logger4repl2spl "test@g_mygroup:informix.t1" "select * from t1" "test@g_loopback:informix.t1" "select * from t1"
```

```
$ cdr start repl rep1
```

Use case 1 – Build staging table for data changes (--jsonsplname example)

```
create database test with log;
```

```
create table t1 (c1 int , c2 int);
```

```
create table t2 (col1 int , col2 float);
```

```
create table staging (data json);
```

```
create procedure logger4repl2spl (data json)
```

```
    insert into staging values (data);
```

```
end procedure;
```

```
$ cdr define repl rep1 -C always -S row -M g_cdr_utm_nag_1 -A -R --jsonsplname=logger4repl2spl "test@g_mygroup:informix.t1" "select *  
from t1" "test@g_loopback:informix.t1" "select * from t1"
```

```
$ cdr define repl rep2 -C always -S row -M g_cdr_utm_nag_1 -A -R --jsonsplname=logger4repl2spl "test@g_mygroup:informix.t2" "select *  
from t2" "test@g_loopback:informix.t2" "select * from t2"
```

```
$ cdr start repl rep1
```

```
$ cdr start repl rep2
```

Use case 2 – Realtime aggregation framework

```
drop database retaildb;
```

```
create database retaildb with log;
```

```
create table sales (customerid int, storeid int , bill_amount float);
```

```
create table sales_summary(storeid int , s_count int, s_sum float, s_avg float, s_min float, s_max float );
```

```
CREATE PROCEDURE store_agg(opType char(1), srcid integer, committime integer, txnid bigint, customerid_bef integer, storeid_bef int, bill_amount_bef float, customerid int, storeid_aft int , bill_amount float)
```

```
-----
```


```
-----
```

```
END PROCEDURE;
```

```
$ cdr define repl rep1 -C always -S row -M utm_group_1 -A -R --serial --splname=store_agg "retaildb@g_mygroup:informix.sales" "select * from sales" "retaildb@g_loopback:informix.sales" "select * from sales"
```

```
$ cdr start repl rep1
```

Use case 3: Leader board calculation



Use case 4 – Publish data to MQTT using Java or C UDR

```
create database mqtt with log;
```

```
create table customer (name char(128), id int);
```

```
execute procedure sqlj.install_jar('file:$INFORMIXDIR/jars/mqtt_trigger.jar', 'mqtt_trigger.jar',1);
```

```
$cdr define repl mqrepl -C always -S row -M g_informix -A -R --serial --jsonsplname=mqtt_put  
"mqtt@g_informix:informix.customer" "select * from customer" "mqtt@g_lb:informix.customer" "select * from customer"
```

```
$ dbaccess mqtt -
```

```
> insert into customer values("Bill", 1);
```

```
$ mosquitto_sub -t 'test/topic' -v
```

```
test/topic
```

```
{"operation":"insert","table":"customer","owner":"informix","database":"mqtt","txnid":21475983636,"commit_time":1550879224,"rowdata":{"name":"Bill","id":1 }}
```

'cdr list repl' sample output

\$cdr list repl

CURRENTLY DEFINED REPLICATES

REPLICATE: rep2

STATE: Active ON:utm_group_1

CONFLICT: Always Apply

FREQUENCY: immediate

QUEUE SIZE: 0

PARTICIPANT: retaildb:informix.sales

OPTIONS: row,ris,ats,fullrow,loopback,rep12spl

REPLID: 65539 / 0x10003

REPLMODE: PRIMARY ON:utm_group_1

APPLY-AS: INFORMIX ON:utm_group_1

REPLTYPE: Master

Demo